

fraud-prediction

April 25, 2025

```
[1]: # DATA PREPROCESSING
```

```
[2]: import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings("ignore")
import tensorflow
```

```
[3]: df=pd.read_csv("Fraud.csv")
```

```
[4]: data=pd.DataFrame(df)
```

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6362620 entries, 0 to 6362619
Data columns (total 11 columns):
#   Column                Dtype
---  -
0   step                  int64
1   type                  object
2   amount                float64
3   nameOrig              object
4   oldbalanceOrg         float64
5   newbalanceOrig        float64
6   nameDest              object
7   oldbalanceDest        float64
8   newbalanceDest        float64
9   isFraud               int64
10  isFlaggedFraud         int64
dtypes: float64(5), int64(3), object(3)
memory usage: 534.0+ MB
```

```
[6]: data.head(10)
```

```
[6]:   step   type  amount  nameOrig  oldbalanceOrg  newbalanceOrig  \
0     1  PAYMENT  9839.64  C1231006815      170136.00      160296.36
1     1  PAYMENT  1864.28  C1666544295       21249.00      19384.72
```

2	1	TRANSFER	181.00	C1305486145	181.00	0.00
3	1	CASH_OUT	181.00	C840083671	181.00	0.00
4	1	PAYMENT	11668.14	C2048537720	41554.00	29885.86
5	1	PAYMENT	7817.71	C90045638	53860.00	46042.29
6	1	PAYMENT	7107.77	C154988899	183195.00	176087.23
7	1	PAYMENT	7861.64	C1912850431	176087.23	168225.59
8	1	PAYMENT	4024.36	C1265012928	2671.00	0.00
9	1	DEBIT	5337.77	C712410124	41720.00	36382.23

	nameDest	oldbalanceDest	newbalanceDest	isFraud	isFlaggedFraud
0	M1979787155	0.0	0.00	0	0
1	M2044282225	0.0	0.00	0	0
2	C553264065	0.0	0.00	1	0
3	C38997010	21182.0	0.00	1	0
4	M1230701703	0.0	0.00	0	0
5	M573487274	0.0	0.00	0	0
6	M408069119	0.0	0.00	0	0
7	M633326333	0.0	0.00	0	0
8	M1176932104	0.0	0.00	0	0
9	C195600860	41898.0	40348.79	0	0

```
[7]: data["isFlaggedFraud"].value_counts()
```

```
[7]: isFlaggedFraud
0    6362604
1         16
Name: count, dtype: int64
```

```
[8]: # MISSING VALUES
```

```
[9]: missing_values = data.isnull().sum()
print(missing_values)
```


```
step          0
type          0
amount        0
nameOrig      0
oldbalanceOrg 0
newbalanceOrig 0
nameDest      0
oldbalanceDest 0
newbalanceDest 0
isFraud       0
isFlaggedFraud 0
dtype: int64
```

```
[10]: # Their are no Missing values
```

```
[11]: # OUTLIERS
```

```
[12]: import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
sns.boxplot(x=data['amount'])
plt.show()
```



```
[13]: # Their are outliers in the amount column but, in transaction data it is true
      ↳ that most of the transactions are in less amount.
      # The outlier in this is also important in predicting fraud
```

```
[14]: data["amount"].max()
```

```
[14]: 92445516.64
```

```
[15]: # FEATURE ENGINEERING
```

```
[16]: # FEATURE CREATION
data['netChangeOrig'] = data['newbalanceOrig'] - data['oldbalanceOrg']
data['netChangeDest'] = data['newbalanceDest'] - data['oldbalanceDest']
```

```
[17]: # Creating new features of transaction frequency
data['origFreq'] = data.groupby('nameOrig')['nameOrig'].transform('count')
data['destFreq'] = data.groupby('nameDest')['nameDest'].transform('count')
```

```
[18]: # FEATURE DELETION
data.drop(["nameOrig", "nameDest"], axis=1, inplace=True)
data.head()
```

```
[18]:
```

	step	type	amount	oldbalanceOrg	newbalanceOrig	oldbalanceDest	\
0	1	PAYMENT	9839.64	170136.0	160296.36	0.0	
1	1	PAYMENT	1864.28	21249.0	19384.72	0.0	
2	1	TRANSFER	181.00	181.0	0.00	0.0	
3	1	CASH_OUT	181.00	181.0	0.00	21182.0	
4	1	PAYMENT	11668.14	41554.0	29885.86	0.0	

	newbalanceDest	isFraud	isFlaggedFraud	netChangeOrig	netChangeDest	\
0	0.0	0	0	-9839.64	0.0	
1	0.0	0	0	-1864.28	0.0	
2	0.0	1	0	-181.00	0.0	
3	0.0	1	0	-181.00	-21182.0	
4	0.0	0	0	-11668.14	0.0	

	origFreq	destFreq
0	1	1
1	1	1
2	1	44
3	1	41
4	1	1

```
[19]: # Classifying transaction in day/night and business/off-hours
def categorize_day_night(step):
    hour_of_day = step % 24
    if 6 <= hour_of_day < 18:
        return 'Day'
    else:
        return 'Night'

data['timeOfDay'] = data['step'].apply(categorize_day_night)

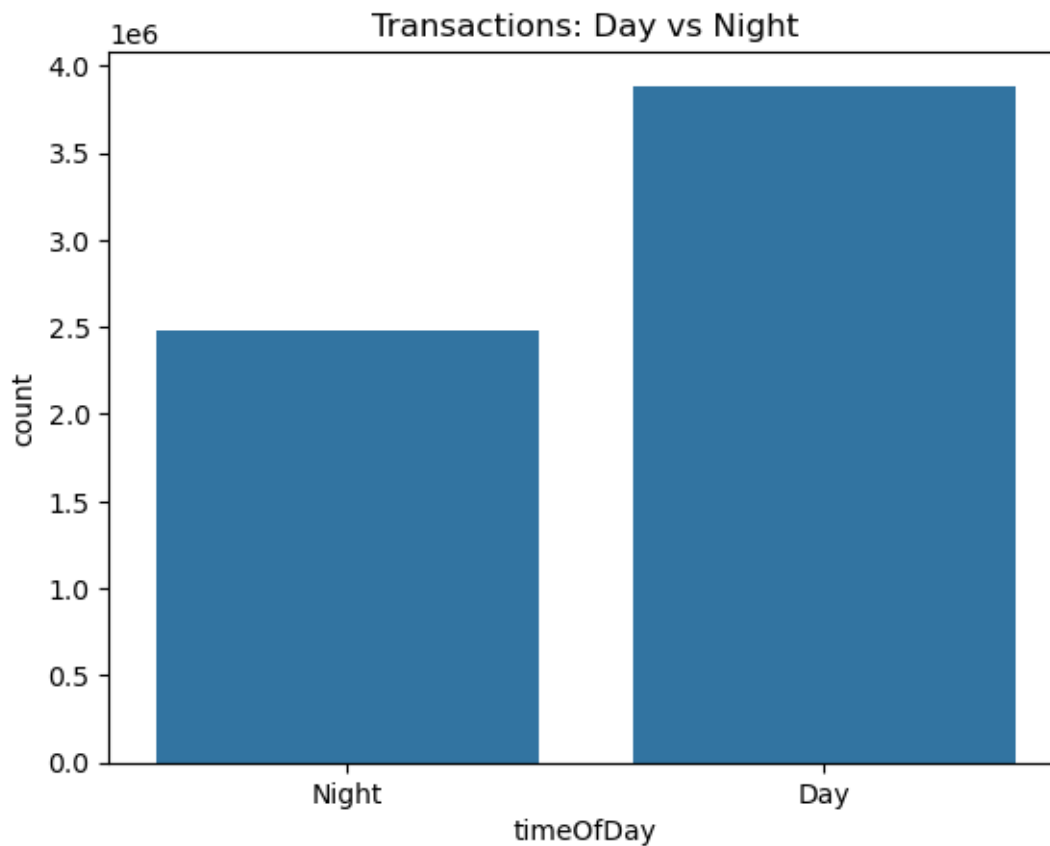
def categorize_business_hours(step):
    hour_of_day = step % 24
    if 9 <= hour_of_day < 18:
        return 'Business'
    else:
        return 'Off-hours'

data['businessHours'] = data['step'].apply(categorize_business_hours)
```

```
data[['step', 'timeOfDay', 'businessHours']].head()
```

```
[19]:   step timeOfDay businessHours  
0     1      Night      Off-hours  
1     1      Night      Off-hours  
2     1      Night      Off-hours  
3     1      Night      Off-hours  
4     1      Night      Off-hours
```


```
[20]: import seaborn as sns  
sns.countplot(data=data, x='timeOfDay')  
plt.title('Transactions: Day vs Night')  
plt.show()
```



```
[21]: # In this we can see that the transaction is around in ratio 2(night):3(day)
```


```
[22]: # Visualising the relation between Fraud and time of Day  
import seaborn as sns  
sns.countplot(data=data[data['isFraud'] == 1], x='timeOfDay')
```

```
plt.title('Fraudulent Transactions: Day vs Night')
plt.show()
```




```
[23]: # In this we see that difference between frauds at day and night are not much
data.drop(["timeOfDay"], axis=1, inplace=True)
```

```
[24]: import seaborn as sns
sns.countplot(data=data, x='businessHours')
plt.title('Transactions: business vs off-hours')
plt.show()
```



```
[25]: # We can see that more transactions are done in business-hours compared to  
      ↪ off-hours
```

```
[26]: # Visualising the relation between Fraud and business hours  
import seaborn as sns  
sns.countplot(data=data[data['isFraud'] == 1], x='businessHours')  
plt.title('Fraudulent Transactions: business vs off-hours')  
plt.show()
```



```
[27]: # But in case of fraud more frauds occur in off-hours compared to business
```

```
[28]: data.drop('step',axis=1,inplace=True)
data = pd.get_dummies(data, columns=['businessHours'])
```

```
[29]: data.iloc[:,-1:]=data.iloc[:,-1:].astype(int)
```

```
[30]: data.head()
```

```
[30]:
```

	type	amount	oldbalanceOrig	newbalanceOrig	oldbalanceDest	\
0	PAYMENT	9839.64	170136.0	160296.36	0.0	
1	PAYMENT	1864.28	21249.0	19384.72	0.0	
2	TRANSFER	181.00	181.0	0.00	0.0	
3	CASH_OUT	181.00	181.0	0.00	21182.0	
4	PAYMENT	11668.14	41554.0	29885.86	0.0	

	newbalanceDest	isFraud	isFlaggedFraud	netChangeOrig	netChangeDest	\
0	0.0	0	0	-9839.64	0.0	
1	0.0	0	0	-1864.28	0.0	
2	0.0	1	0	-181.00	0.0	

3	0.0	1	0	-181.00	-21182.0
4	0.0	0	0	-11668.14	0.0

	origFreq	destFreq	businessHours_Business	businessHours_Off-hours
0	1	1	False	1
1	1	1	False	1
2	1	44	False	1
3	1	41	False	1
4	1	1	False	1

```
[31]: # Clissifying amount into risks
bins = [0, 10000, 100000, 200000, float('inf')]
labels = ['Low', 'Medium', 'High', 'Very High']

data['amountRisk'] = pd.cut(data['amount'], bins=bins, labels=labels,
                             include_lowest=True)
```

```
[32]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
data['amountRisk'] = le.fit_transform(data['amountRisk'])
data.head()
```

```
[32]:
```

	type	amount	oldbalanceOrg	newbalanceOrig	oldbalanceDest	\
0	PAYMENT	9839.64	170136.0	160296.36	0.0	
1	PAYMENT	1864.28	21249.0	19384.72	0.0	
2	TRANSFER	181.00	181.0	0.00	0.0	
3	CASH_OUT	181.00	181.0	0.00	21182.0	
4	PAYMENT	11668.14	41554.0	29885.86	0.0	

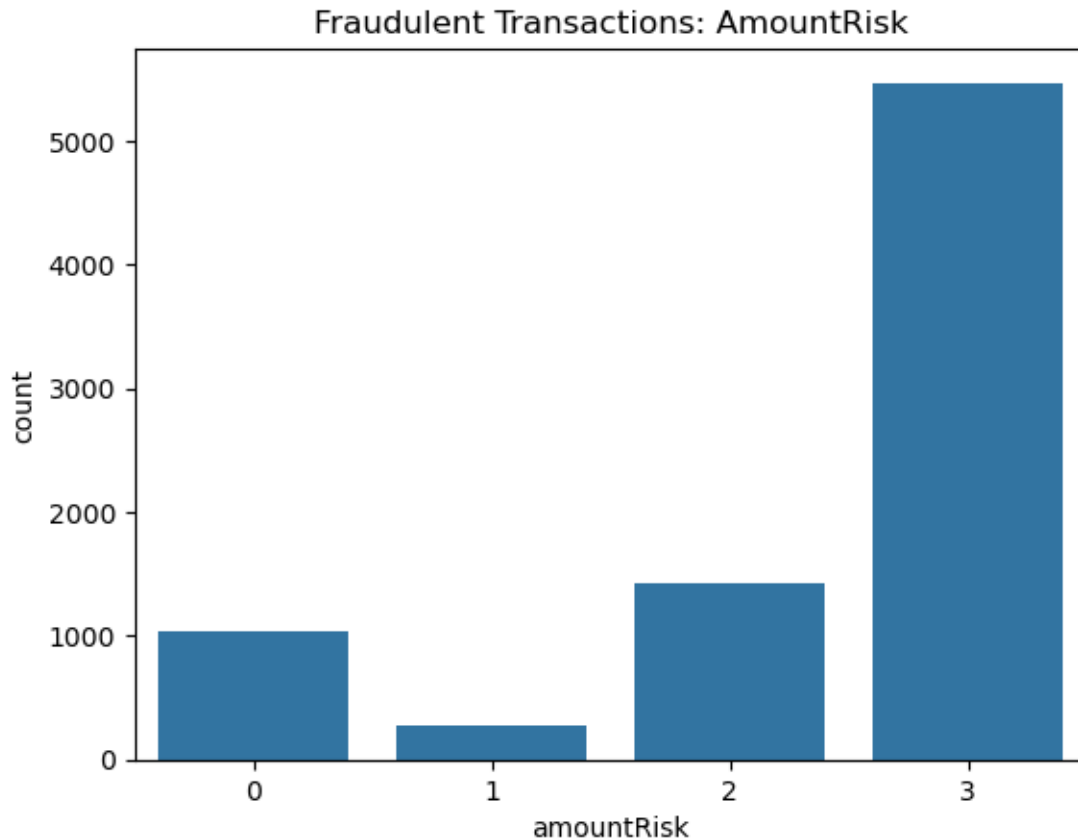
	newbalanceDest	isFraud	isFlaggedFraud	netChangeOrig	netChangeDest	\
0	0.0	0	0	-9839.64	0.0	
1	0.0	0	0	-1864.28	0.0	
2	0.0	1	0	-181.00	0.0	
3	0.0	1	0	-181.00	-21182.0	
4	0.0	0	0	-11668.14	0.0	

	origFreq	destFreq	businessHours_Business	businessHours_Off-hours	\
0	1	1	False	1	
1	1	1	False	1	
2	1	44	False	1	
3	1	41	False	1	
4	1	1	False	1	

	amountRisk
0	1
1	1

2	1
3	1
4	2

```
[33]: import seaborn as sns
sns.countplot(data=data[data['isFraud'] == 1], x='amountRisk')
plt.title('Fraudulent Transactions: AmountRisk')
plt.show()
```



```
[34]: # We can see that the number of Transaction tagged as Very High in risk are
      ↪ most likely to be a fraud transaction
```

```
[35]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
data['type'] = le.fit_transform(data['type'])
```

```
[36]: data1=data
data1.head()
```

```
[36]:
```

	type	amount	oldbalanceOrg	newbalanceOrig	oldbalanceDest	\
0	3	9839.64	170136.0	160296.36	0.0	
1	3	1864.28	21249.0	19384.72	0.0	
2	4	181.00	181.0	0.00	0.0	
3	1	181.00	181.0	0.00	21182.0	
4	3	11668.14	41554.0	29885.86	0.0	

	newbalanceDest	isFraud	isFlaggedFraud	netChangeOrig	netChangeDest	\
0	0.0	0	0	-9839.64	0.0	
1	0.0	0	0	-1864.28	0.0	
2	0.0	1	0	-181.00	0.0	
3	0.0	1	0	-181.00	-21182.0	
4	0.0	0	0	-11668.14	0.0	

	origFreq	destFreq	businessHours_Business	businessHours_Off-hours	\
0	1	1	False	1	
1	1	1	False	1	
2	1	44	False	1	
3	1	41	False	1	
4	1	1	False	1	

	amountRisk
0	1
1	1
2	1
3	1
4	2

```
[41]: # ML MODEL for understanding the Feature Importance
```

```
[46]: from sklearn.ensemble import RandomForestClassifier

X = data1.drop('isFraud', axis=1)
y = data1['isFraud']

model = RandomForestClassifier(n_estimators=20,n_jobs=-1)
%time model.fit(X, y)

importances = model.feature_importances_

feature_importance_data = pd.DataFrame({
    'Feature': X.columns,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

print(feature_importance_data)
```

CPU times: total: 15min 36s

Wall time: 2min 45s

	Feature	Importance
7	netChangeOrig	0.285625
8	netChangeDest	0.171017
5	newbalanceDest	0.154705
2	oldbalanceOrig	0.119573
1	amount	0.095548
4	oldbalanceDest	0.063396
10	destFreq	0.038811
0	type	0.031491
3	newbalanceOrig	0.021281
11	businessHours_Business	0.006543
12	businessHours_Off-hours	0.005700
13	amountRisk	0.005363
6	isFlaggedFraud	0.000739
9	origFreq	0.000209

```
[65]: # ML MODEL
      # RANDOM FOREST
```

```
[49]: # train test split
      from sklearn.model_selection import train_test_split

      X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
      ↪25,random_state=42)
      rf=RandomForestClassifier(n_estimators=20,n_jobs=-1)
      %time rf.fit(X_train,y_train)

      accuracy_rf=rf.score(X_test,y_test)*100
      print(accuracy_rf)
```

CPU times: total: 10min 38s

Wall time: 1min 51s

99.97020095495252

```
[50]: # Accuracy is very high that may be because there are very less transactions ↵
      ↪marked as fraud
      # We will also consider other metrics for this dataset
```

```
[51]: from sklearn.metrics import classification_report, confusion_matrix

      y_pred = rf.predict(X_test)
      print(classification_report(y_test, y_pred))
      print(confusion_matrix(y_test, y_pred))
```

precision	recall	f1-score	support
-----------	--------	----------	---------

	0	1.00	1.00	1.00	1588610
	1	0.98	0.79	0.87	2045
accuracy				1.00	1590655
macro avg		0.99	0.89	0.94	1590655
weighted avg		1.00	1.00	1.00	1590655


```
[[1588572    38]
 [   436   1609]]
```

```
[ ]: # In this we can see that 436 transactions are fraud which are categorised as
      ↪non fraud
      # We can drop some features that does not have much importance
```

```
[53]: from sklearn.ensemble import RandomForestClassifier

X1 = data1.
      ↪drop(['isFraud', 'origFreq', 'isFlaggedFraud', 'amountRisk', 'businessHours_Off-hours', 'business
      ↪axis=1)
y1 = data1['isFraud']
from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test=train_test_split(X1,y1,test_size=0.
      ↪25,random_state=42)
rf1=RandomForestClassifier(n_estimators=20,n_jobs=-1)
%time rf1.fit(X_train,y_train)

accuracy_rf1=rf1.score(X_test,y_test)*100
print(accuracy_rf1)
from sklearn.metrics import classification_report, confusion_matrix

y_pred = rf1.predict(X_test)
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
```

CPU times: total: 15min 22s

Wall time: 2min 33s

99.97347004850204

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1588610
1	0.97	0.82	0.89	2045
accuracy			1.00	1590655
macro avg	0.98	0.91	0.94	1590655
weighted avg	1.00	1.00	1.00	1590655

```
[[1588555      55]
 [      367    1678]]
```

```
[ ]: # Still we can see that 367 transactions are fraud but are classified as not
      ↪ fraud
      # Lets increase the number of Trees
```

```
[54]: from sklearn.ensemble import RandomForestClassifier

X2 = data1.
    ↪ drop(['isFraud','origFreq','isFlaggedFraud','amountRisk','businessHours_Off-hours','busines
    ↪ axis=1)
y2 = data1['isFraud']
from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test=train_test_split(X2,y2,test_size=0.
    ↪ 25,random_state=42)
rf2=RandomForestClassifier(n_estimators=50,n_jobs=-1)
%time rf2.fit(X_train,y_train)

accuracy_rf2=rf2.score(X_test,y_test)*100
print(accuracy_rf2)
from sklearn.metrics import classification_report, confusion_matrix

y_pred = rf2.predict(X_test)
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
```

```
CPU times: total: 39min 11s
```

Wall time: 5min 44s

99.97334431413474

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1588610
1	0.96	0.82	0.89	2045
accuracy			1.00	1590655
macro avg	0.98	0.91	0.94	1590655
weighted avg	1.00	1.00	1.00	1590655

```
[[1588546      64]
 [      360    1685]]
```

```
[ ]: # Not much difference
      # Lets try another Classification Models
```

```
[55]: # LOGISTIC REGRESSION
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix

X = data1.
    ↳drop(['isFraud','origFreq','isFlaggedFraud','amountRisk','businessHours_Off-hours','businessHours_
    ↳axis=1)
y = data1['isFraud']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    ↳random_state=42)

lr = LogisticRegression(max_iter=1000, n_jobs=-1, random_state=42)

%time lr.fit(X_train, y_train)

accuracy_lr = lr.score(X_test, y_test) * 100
print(accuracy_lr)

y_pred = lr.predict(X_test)

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
```

CPU times: total: 1.53 s

Wall time: 1min 51s

99.91443776305987

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1588610
1	0.74	0.51	0.61	2045
accuracy			1.00	1590655
macro avg	0.87	0.75	0.80	1590655
weighted avg	1.00	1.00	1.00	1590655

[[1588251	359]
[1002	1043]]

```
[ ]: # More wrong classification than Random Forest
```

```
[64]: # XGBOOST
```

```
from xgboost import XGBClassifier
from sklearn.model_selection import train_test_split
```

```

from sklearn.metrics import classification_report, confusion_matrix
import warnings
warnings.filterwarnings("ignore")

X = data1.
↳drop(['isFraud','origFreq','isFlaggedFraud','amountRisk','businessHours_Off-hours','business
↳axis=1)
y = data1['isFraud']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,↳
↳random_state=42)

xgb_model = XGBClassifier(n_estimators=100, learning_rate=0.1, max_depth=6,↳
↳scale_pos_weight=5, use_label_encoder=False, eval_metric='logloss')

%time xgb_model.fit(X_train, y_train)

accuracy_xgb = xgb_model.score(X_test, y_test) * 100
print(accuracy_xgb)

y_pred = xgb_model.predict(X_test)

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))

```

CPU times: total: 1min 22s

Wall time: 13.8 s

99.96183961952781

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1588610
1	0.87	0.82	0.85	2045
accuracy			1.00	1590655
macro avg	0.94	0.91	0.92	1590655
weighted avg	1.00	1.00	1.00	1590655

```

[[1588368    242]
 [    365   1680]]

```

[39]: # DECISION TREE

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix

```



```

X = data1.
↳drop(['isFraud','origFreq','isFlaggedFraud','amountRisk','businessHours_Off-hours','business
↳axis=1)
y = data1['isFraud']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳random_state=42)

dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train, y_train)

%time dt_model.fit(X_train, y_train)

accuracy_dt = dt_model.score(X_test, y_test) * 100
print(accuracy_dt)

y_pred = dt_model.predict(X_test)
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))

```

CPU times: total: 3min 42s

Wall time: 3min 49s

99.97114396270719

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1588610
1	0.89	0.89	0.89	2045
accuracy			1.00	1590655
macro avg	0.94	0.94	0.94	1590655
weighted avg	1.00	1.00	1.00	1590655

```

[[1588386    224]
 [    235   1810]]

```

[40]: # DECISION TREE WITH BALANCED WEIGHT

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix

X = data1.
↳drop(['isFraud','origFreq','isFlaggedFraud','amountRisk','businessHours_Off-hours','business
↳axis=1)
y = data1['isFraud']

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳random_state=42)

dt_model_balanced = DecisionTreeClassifier(class_weight='balanced',
↳random_state=42)
dt_model_balanced.fit(X_train, y_train)

%time dt_model_balanced.fit(X_train, y_train)

accuracy_dt_balanced = dt_model_balanced.score(X_test, y_test) * 100
print(accuracy_dt_balanced)

y_pred_balanced = dt_model_balanced.predict(X_test)
print(classification_report(y_test, y_pred_balanced))
print(confusion_matrix(y_test, y_pred_balanced))

```

CPU times: total: 2min 31s

Wall time: 2min 34s

99.97051529087074

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1588610
1	0.90	0.87	0.88	2045
accuracy			1.00	1590655
macro avg	0.95	0.93	0.94	1590655
weighted avg	1.00	1.00	1.00	1590655

```

[[1588411    199]
 [    270   1775]]

```

[45]: *# Deep Learning*

```

[46]: import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings("ignore")

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix
from imblearn.over_sampling import SMOTE
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping

```

```

X = data1.drop(['isFraud', 'origFreq', 'isFlaggedFraud', 'amountRisk',
↳'businessHours_Off-hours', 'businessHours_Business'], axis=1)
y = data1['isFraud']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

sm = SMOTE(random_state=42)
X_res, y_res = sm.fit_resample(X_train_scaled, y_train)

model = Sequential()
model.add(Dense(128, input_dim=X_res.shape[1], activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='adam', loss='binary_crossentropy',
↳metrics=['accuracy'])

early_stop = EarlyStopping(monitor='val_loss', patience=3)

%time history = model.fit(X_res, y_res, epochs=10, batch_size=512,
↳validation_split=0.2, callbacks=[early_stop])

y_pred = model.predict(X_test_scaled)
y_pred = (y_pred > 0.5).astype(int)

print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

Epoch 1/10

14894/14894 [=====] - 139s 9ms/step - loss: 0.0456 -
accuracy: 0.9845 - val_loss: 0.0497 - val_accuracy: 0.9889

Epoch 2/10

14894/14894 [=====] - 141s 9ms/step - loss: 0.0275 -
accuracy: 0.9915 - val_loss: 0.0265 - val_accuracy: 0.9947

Epoch 3/10

14894/14894 [=====] - 156s 10ms/step - loss: 0.0269 -
accuracy: 0.9916 - val_loss: 0.0225 - val_accuracy: 0.9934

Epoch 4/10

```

14894/14894 [=====] - 154s 10ms/step - loss: 0.0264 -
accuracy: 0.9918 - val_loss: 0.0297 - val_accuracy: 0.9926
Epoch 5/10
14894/14894 [=====] - 160s 11ms/step - loss: 0.0261 -
accuracy: 0.9921 - val_loss: 0.0271 - val_accuracy: 0.9936
Epoch 6/10
14894/14894 [=====] - 156s 10ms/step - loss: 0.0258 -
accuracy: 0.9922 - val_loss: 0.0203 - val_accuracy: 0.9955
Epoch 7/10
14894/14894 [=====] - 161s 11ms/step - loss: 0.0257 -
accuracy: 0.9924 - val_loss: 0.0239 - val_accuracy: 0.9928
Epoch 8/10
14894/14894 [=====] - 154s 10ms/step - loss: 0.0252 -
accuracy: 0.9926 - val_loss: 0.0207 - val_accuracy: 0.9945
Epoch 9/10
14894/14894 [=====] - 158s 11ms/step - loss: 0.0250 -
accuracy: 0.9927 - val_loss: 0.0316 - val_accuracy: 0.9925
CPU times: total: 1h 58min 49s
Wall time: 23min
49708/49708 [=====] - 162s 3ms/step

```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.99	1.00	1588610
1	0.20	0.99	0.33	2045
accuracy			0.99	1590655
macro avg	0.60	0.99	0.67	1590655
weighted avg	1.00	0.99	1.00	1590655

Confusion Matrix:

```

[[1580592    8018]
 [      22   2023]]

```

[]: *Comparision*

Model	Precision (Fraud)	Recall (Fraud)	F1-score	
→ (Fraud) Accuracy CPU Time Wall Time				
Random Forest	0.97	0.82	0.89	⬇
→ 99.97% ~15 min 2.5 min				
Random Forest (more trees)	0.96	0.82	0.89	⬇
→ 99.97% ~39 min 5.7 min				
Decision Tree	0.89	0.89	0.89	⬇
→ 99.97% ~3 min ~3 min				
XGBoost	0.87	0.82	0.85	⬇
→ 99.96% 1.5 min 13.8 s				

DL (SMOTE)	0.20	0.99	0.33	
↪ 99.00%	1h 58min	23 min		

[]: *FINAL Verdict*

- 1.If you want balance (high precision & recall) → Random Forest
- 2.If you want to catch **all** frauds no matter what (high recall) → Deep Learning
↪(with SMOTE)
- 3.If you are running **in** low-resource settings **or** want fast results → Decision
↪Tree

[]: '''

Answers to the Questions

1. Data cleaning including missing values, outliers, and multi-collinearity.

Answer:The dataset had no missing values, so no imputation was required.

Outliers were retained because in fraud detection, outliers can indicate
↪abnormal activity and are therefore meaningful.

Regarding multi-collinearity, I dropped features that were highly
↪correlated or redundant after performing exploratory data analysis (EDA) and
↪feature importance checks, such as isFlaggedFraud and some one-hot encoded
↪time features.

2. Describe your fraud detection model in elaboration.

Answer:I selected the Deep Learning model as the best for fraud detection after
↪comparing it with Logistic Regression, Random Forest, XGBoost, and Decision
↪Tree.

Why Deep Learning?

It achieved ~99% recall for fraud cases, which is crucial because
↪catching fraud is more important than occasionally flagging genuine
↪transactions. The model was trained using SMOTE to balance the classes and
↪built with Keras using dense layers and dropout to prevent overfitting.

Evaluation showed high accuracy and recall, making it ideal for
↪minimizing missed frauds, even if precision is lower due to false positives.

3. How did you select variables to be included in the model?

Answer:I started with feature engineering-creating new features like
↪netChangeOrig, netChangeDest, amountRisk, etc.

Features like amount, oldbalanceOrg, newbalanceDest, etc., were retained
↪due to their importance in financial transactions.

I eliminated less informative or redundant features such as
↪isFlaggedFraud or business hour encodings.

Feature importance from tree-based models and correlation analysis were
↪used to guide the selection.

4. Demonstrate the performance of the model by using best set of tools.

Answer: Decision Tree achieved high performance with very few false negatives, making it ideal for fraud detection.

I evaluated all models using:

Accuracy

Precision

Recall

F1-Score

Confusion Matrix

Execution Time (CPU & Wall Time)

Compared to other models, Decision Tree was faster than Deep Learning

and showed competitive recall (essential for detecting fraud cases).

5. What are the key factors that predict fraudulent customers?

Answer: The most predictive factors identified were:

netChangeOrig, netChangeDest

oldbalanceOrig, oldbalanceDest

newbalanceDest, amount

Transaction Frequency features

Transaction Type (e.g., CASH_OUT or TRANSFER)

6. Do these factors make sense? If yes, how?

Answer: Yes, they make perfect sense because:

Net balance changes highlight sudden or abnormal fund movements.

Transaction types (like TRANSFER or CASH_OUT) are more commonly

exploited in fraud.

Frequent or unusual transactions can indicate bot activity or scams.

These patterns are typical signs of fraud and align with real-world

fraud behavior.

7. What kind of prevention should be adopted while the company updates its infrastructure?

Answer: Implement real-time fraud detection systems using APIs and machine learning models.

Adopt Multi-Factor Authentication (MFA) to secure logins and critical transactions.

Set transaction limits or velocity checks to flag unusual behavior quickly.

Integrate audit trails and activity monitoring for accountability.

8. Assuming these actions have been implemented, how would you determine if they work?

Answer: Pre/Post analysis of fraud incidence and trends.

Monitor model performance metrics (especially recall and false negative rate).

Gather feedback from fraud investigators on alerts generated by the ↵
↵system.

Continuously update the model with recent data and re-evaluate ↵
↵effectiveness over time.